

Name 1:

Name 2:

Suppose that each heap block can hold 10 tuples, and each index block can hold 100 index records (for B+tree indexes). Nodes of the index are 70% full. The relation contains 1 million records.

Determine, for each of the structures determine:


- The total number of blocks needed by the heap and the leaves of the index.
- The average number of disk I/Os needed to retrieve a given search key.

Assume that nothing is in memory initially, and that the search key is the primary key of the table.

- The table is an unsorted heap, packed 10 to a block. The B+tree is dense.
- The table is a sorted heap with 10 records per block. The B+tree is sparse and only the first record in each block is in the index.

a) Table : $\#blocks = \frac{\#records}{\#records\ per\ page} = \frac{10^6}{10} = 10^5$

Index :



Index is dense
so one index record per
tuple.

$\Rightarrow \frac{10^6}{70}$ blocks in
leaves of index

the height of the tree is $\lceil \log_{70}(10^6) \rceil$

$$h = \left\lceil \frac{\log_{10}(10^6)}{\log_{10}(70)} \right\rceil = \left\lceil \frac{6}{1.8} \right\rceil = 4$$

How expensive:

- Using index, $h + 1$.

- Using seq scan

10^5 blocks if not found:

$\frac{10^5}{2}$ block, average if found.

(because it is a primary key)

10^5 worst case

b) Sparse Index.

Index will only contain one record per block.

We need $\frac{10^5}{70}$ blocks in the index. Heap is 10^5 also.

$$h = \left\lceil \frac{10^5}{70} \right\rceil = \left\lceil \frac{5}{1.8} \right\rceil = 3$$

So we need to read $h + 1 = 4$ blocks

whether the heap contains the value or not